

Computer Vision

CS308

Feng Zheng

SUSTech CS Vision Intelligence and Perception
Week 4



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY



Content

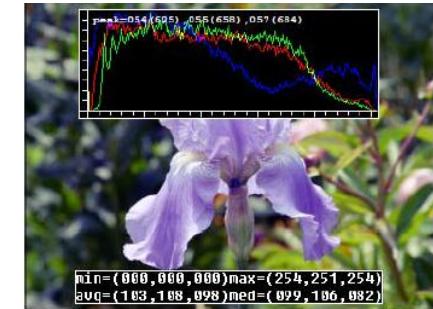
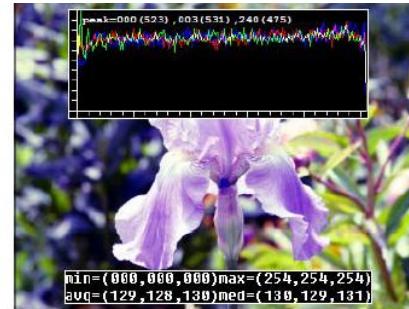
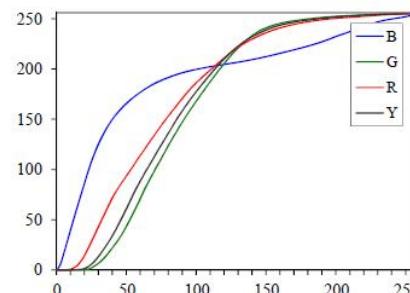
- Brief Review
- Thinking in Frequency
- Pyramids
- Geometric Transformations

Brief Review



Review

- Point Operators



- Linear Filtering

➤ Linear filtering is dot product at each position

- ✓ Not a matrix multiplication
- ✓ Can smooth, sharpen, translate (among many other uses)

➤ Be aware of details

- ✓ Filter size
- ✓ Stride
- ✓ Padding
- ✓ Values to be fixed or learned

Auto ML

45	60	98	127	132	133	137	133
46	65	98	123	126	128	131	133
47	65	96	115	119	123	135	137
47	63	91	107	113	122	138	134
50	59	80	97	110	123	133	134
49	53	68	83	97	113	128	133
50	50	58	70	84	102	116	126
50	50	52	58	69	86	101	120

$$\begin{bmatrix} 0.1 & 0.1 & 0.1 \\ 0.1 & 0.2 & 0.1 \\ 0.1 & 0.1 & 0.1 \end{bmatrix}$$

*

69	95	116	125	129	132
68	92	110	120	126	132
66	86	104	114	124	132
62	78	94	108	120	129
57	69	83	98	112	124
53	60	71	85	100	114

$f(x,y)$

$h(x,y)$

$g(x,y)$



Review

- More Neighborhood Operators



(a)



(b)



(c)



(d)



(e)



(f)

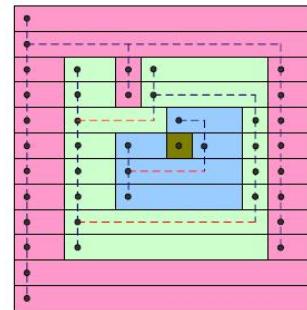
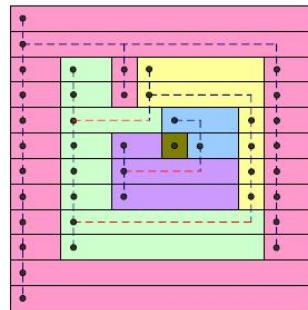
1	2	1	2	4
2	1	3	5	8
1	3	7	6	9
3	4	8	6	7
4	5	7	8	9

(a) median = 4

1	2	1	2	4
2	1	3	5	8
1	3	7	6	9
3	4	8	6	7
4	5	7	8	9

(b) α -mean= 4.6

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0



0	0	0	0	0	1	0	0	0
0	0	1	1	1	0	0	0	0
0	1	1	1	1	1	0	0	0
0	1	1	1	1	1	0	0	0
0	1	1	1	0	0	0	0	0

0	0	0	0	0	1	0	0	0
0	0	1	1	2	0	0	0	0
0	1	2	2	3	1	0	0	0
0	1	2	3	1	1	0	0	0
0	1	2	1	0	0	0	0	0

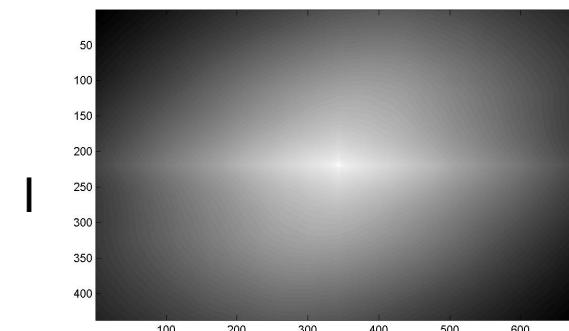
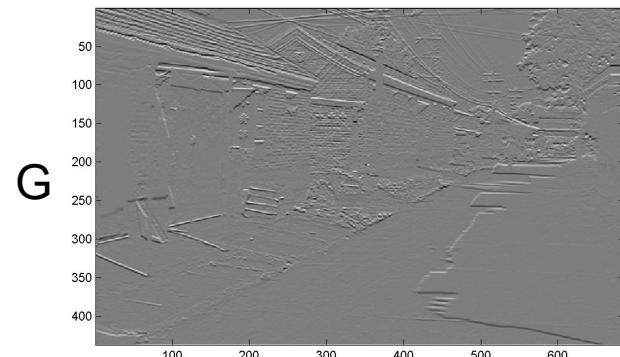
0	0	0	0	0	1	0	0	0
0	0	1	1	2	0	0	0	0
0	1	2	2	3	1	0	0	0
0	1	2	1	1	1	0	0	0
0	1	2	1	0	0	0	0	0

0	0	0	0	0	1	0	0	0
0	0	1	1	1	1	0	0	0
0	1	2	2	2	2	1	0	0
0	1	2	2	2	1	1	0	0
0	1	2	1	1	1	1	0	0

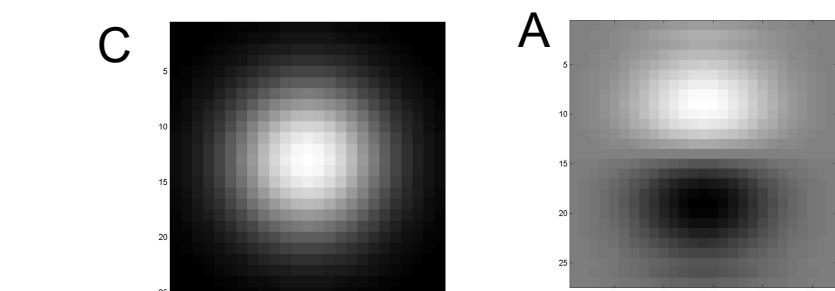
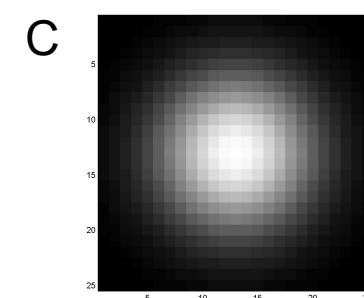


Review

- Fill in the blanks



- Filtering Operator
- a) $() = D * B$
- b) $A = () * ()$
- c) $F = D * ()$
- d) $() = D * D$



Thinking in Frequency

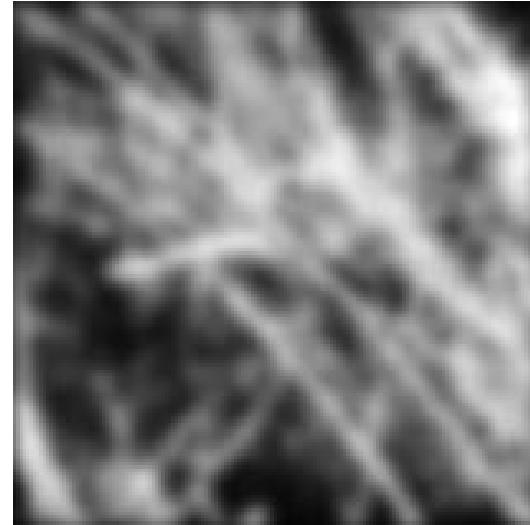


Thinking

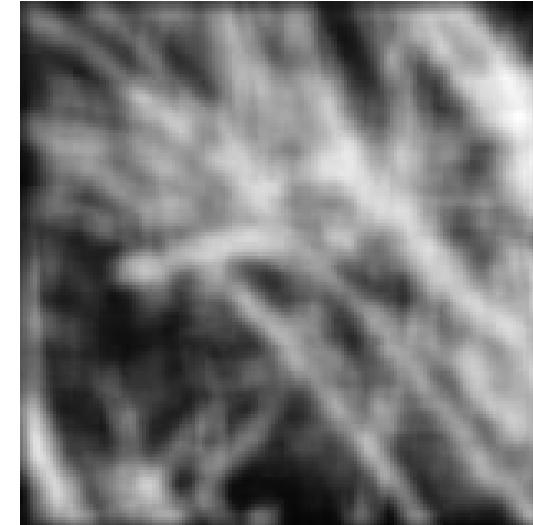
- Why does the Gaussian give a **nice smooth** image, but the square filter give edgy **artifacts**?



Gaussian



Box filter

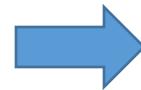


- How is it that a 4MP image can be **compressed** to a few hundred KB without a noticeable change?



Thinking

- Why does a **lower** resolution image still **make sense** to us?
- What do we **lose**?



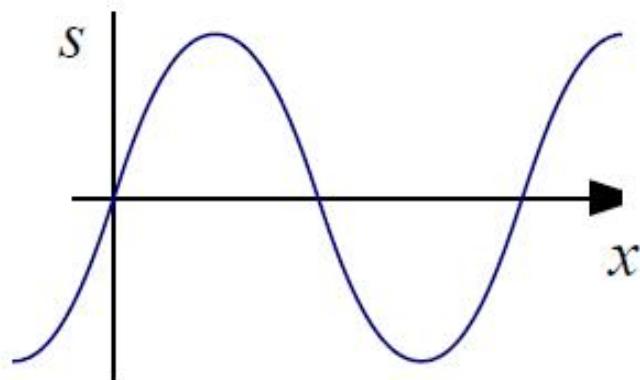


Filtering and Fourier

- How can we analyze what a given filter does to high, medium, and low **frequencies**?
 - Pass a sinusoid of known frequency through the **filter** and observe by how much it is attenuated

$$s(x) = \sin(2\pi f x + \phi_i) = \sin(\omega x + \phi_i)$$

frequency
angular frequency
phase

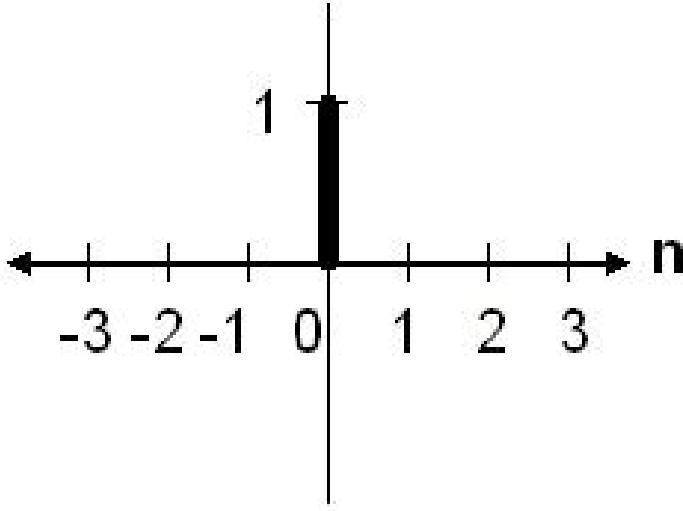




Filtering and Fourier

- Impulse response: $h(x)$

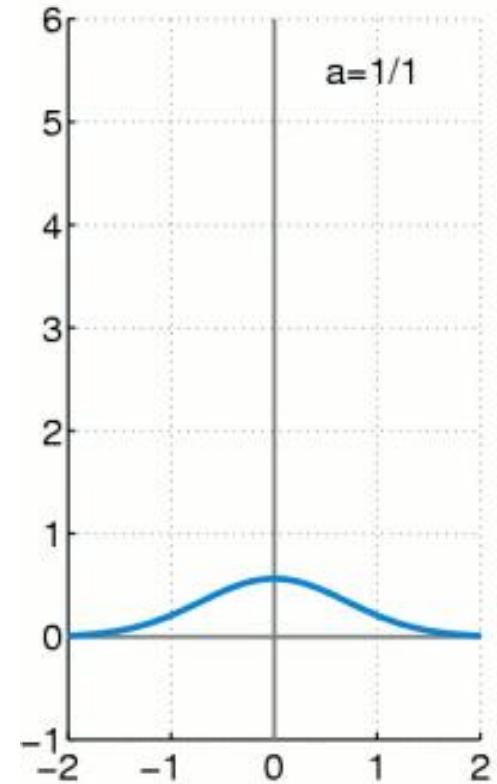
- Continuous function
 - Discrete function



- The impulse function **contains all frequencies**

- If we **convolve** the signal with a filter whose impulse response is $h(x)$, we get another sinusoid of the **same frequency** but different magnitude and phase

$$o(x) = h(x) * s(x) = A \sin(\omega x + \phi_o)$$





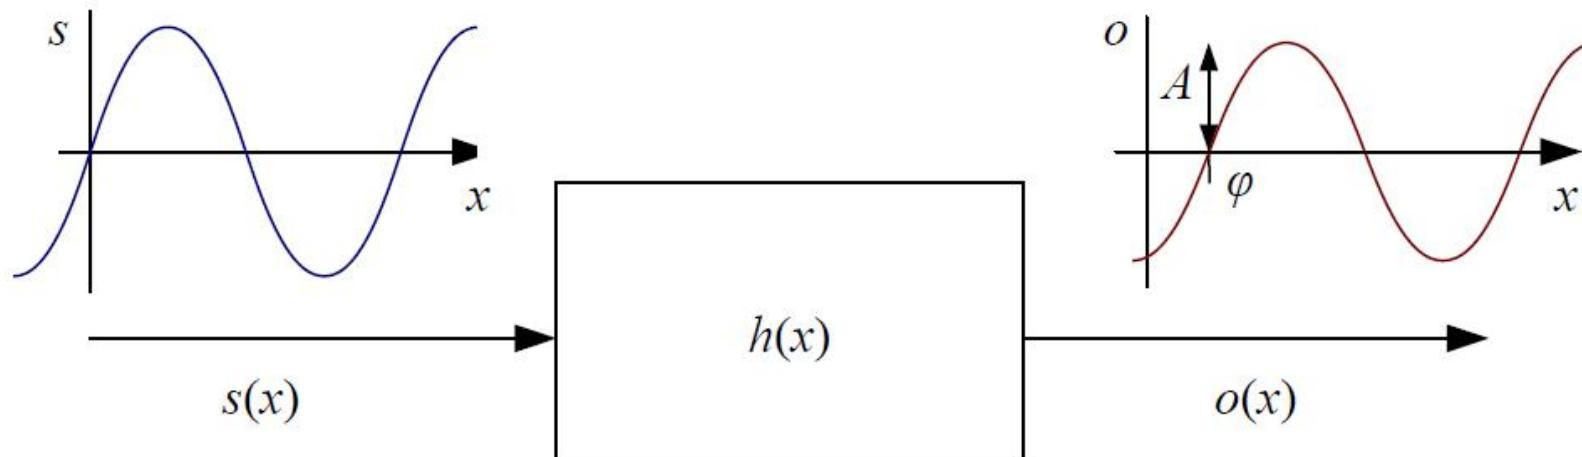
Filtering and Fourier

- Convolution

$$g(\mathbf{x}) = \int f(\mathbf{x} - \mathbf{u})h(\mathbf{u})d\mathbf{u}.$$

- A weighted summation of shifted input signals (sinusoids)
- The summation of a bunch of shifted sinusoids of the same frequency is just a single sinusoid at that frequency (same)

$$o(x) = h(x) * s(x) = A \sin(\omega x + \phi_o)$$





Filtering and Fourier

- The **complex-valued** sinusoid is

$$s(x) = e^{j\omega x} = \cos \omega x + j \sin \omega x$$

- The filtered sinusoid is

$$o(x) = h(x) * s(x) = A e^{j\omega x + \phi}$$

- The Fourier transform in continuous domain

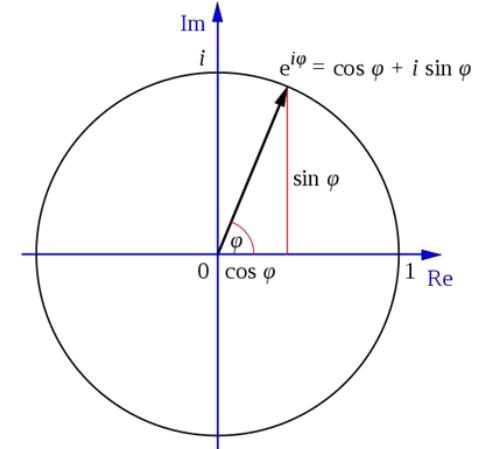
One frequency

$$H(\omega) = \int_{-\infty}^{\infty} h(x) e^{-j\omega x} dx$$

- The transform in discrete domain

$$H(k) = \frac{1}{N} \sum_{x=0}^{N-1} h(x) e^{-j \frac{2\pi k x}{N}}$$

$$e^{ix} = \cos x + i \sin x$$



length of the signal



Filtering and Fourier

- Some useful properties of Fourier transforms

Property	Signal	Transform
superposition	$f_1(x) + f_2(x)$	$F_1(\omega) + F_2(\omega)$
shift	$f(x - x_0)$	$F(\omega)e^{-j\omega x_0}$
reversal	$f(-x)$	$F^*(\omega)$
convolution	$f(x) * h(x)$	$F(\omega)H(\omega)$
correlation	$f(x) \otimes h(x)$	$F(\omega)H^*(\omega)$
multiplication	$f(x)h(x)$	$F(\omega) * H(\omega)$
differentiation	$f'(x)$	$j\omega F(\omega)$
domain scaling	$f(ax)$	$1/aF(\omega/a)$
real images	$f(x) = f^*(x)$	$\Leftrightarrow F(\omega) = F(-\omega)$
Parseval's Theorem	$\sum_x [f(x)]^2$	$= \sum_\omega [F(\omega)]^2$

帕塞瓦尔定理指出，信号的能量在时域和频域中的表示是等价的，意味着信号在时域上的能量等于其傅里叶变换在频域上的能量。



Filtering and Fourier

- Occurring filters and signals

Name	Signal	Transform
impulse	$\delta(x)$	\Leftrightarrow 1
shifted impulse	$\delta(x - u)$	\Leftrightarrow $e^{-j\omega u}$
box filter	$\text{box}(x/a)$	\Leftrightarrow $a\text{sinc}(a\omega)$
tent	$\text{tent}(x/a)$	\Leftrightarrow $a\text{sinc}^2(a\omega)$
Gaussian	$G(x; \sigma)$	\Leftrightarrow $\frac{\sqrt{2\pi}}{\sigma} G(\omega; \sigma^{-1})$
Laplacian of Gaussian	$(\frac{x^2}{\sigma^4} - \frac{1}{\sigma^2})G(x; \sigma)$	\Leftrightarrow $-\frac{\sqrt{2\pi}}{\sigma} \omega^2 G(\omega; \sigma^{-1})$
Gabor	$\cos(\omega_0 x)G(x; \sigma)$	\Leftrightarrow $\frac{\sqrt{2\pi}}{\sigma} G(\omega \pm \omega_0; \sigma^{-1})$
unsharp mask	$(1 + \gamma)\delta(x) - \gamma G(x; \sigma)$	\Leftrightarrow $(1 + \gamma) - \frac{\sqrt{2\pi}\gamma}{\sigma} G(\omega; \sigma^{-1})$
windowed sinc	$r\cos(x/(aW))$ $\text{sinc}(x/a)$	\Leftrightarrow (see Figure 3.29)

Separable kernels

Name	Kernel	Transform	Plot
box-3	$\frac{1}{3} [1 \ 1 \ 1]$	$\frac{1}{3}(1 + 2 \cos \omega)$	
box-5	$\frac{1}{5} [1 \ 1 \ 1 \ 1 \ 1]$	$\frac{1}{5}(1 + 2 \cos \omega + 2 \cos 2\omega)$	
linear	$\frac{1}{4} [1 \ 2 \ 1]$	$\frac{1}{2}(1 + \cos \omega)$	
binomial	$\frac{1}{16} [1 \ 4 \ 6 \ 4 \ 1]$	$\frac{1}{4}(1 + \cos \omega)^2$	
Sobel	$\frac{1}{2} [-1 \ 0 \ 1]$	$\sin \omega$	
corner	$\frac{1}{2} [-1 \ 2 \ -1]$	$\frac{1}{2}(1 - \cos \omega)$	

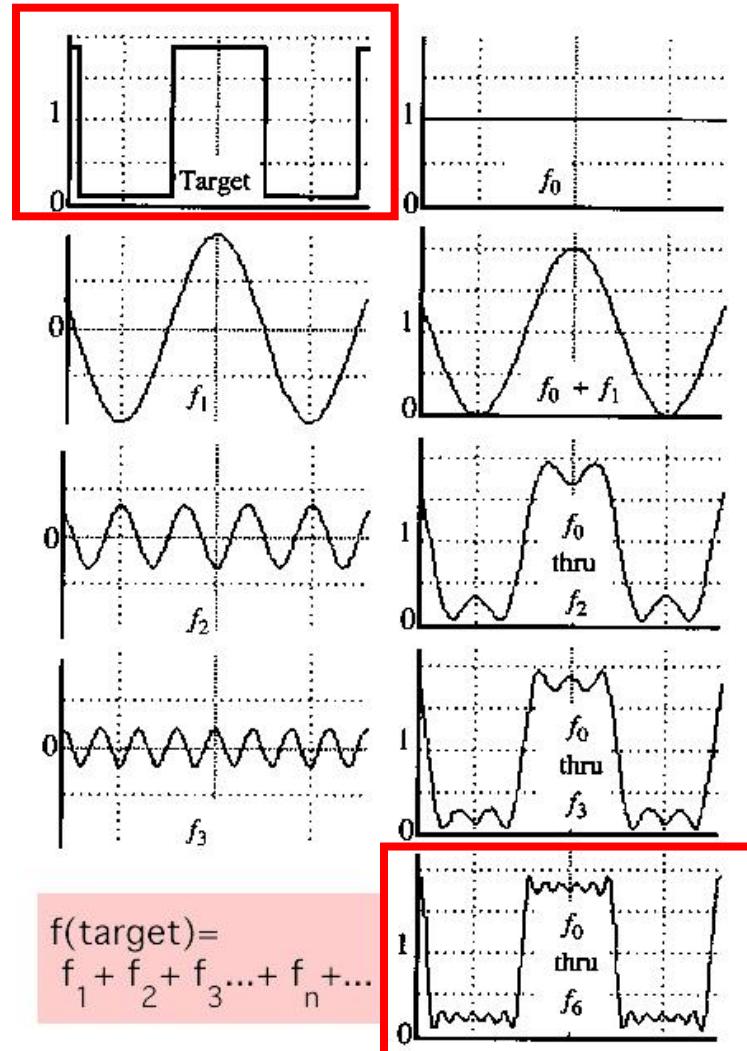


Aexample: A Sum of Sines

- Our building blocks:

$$A \sin(\omega x + \phi)$$

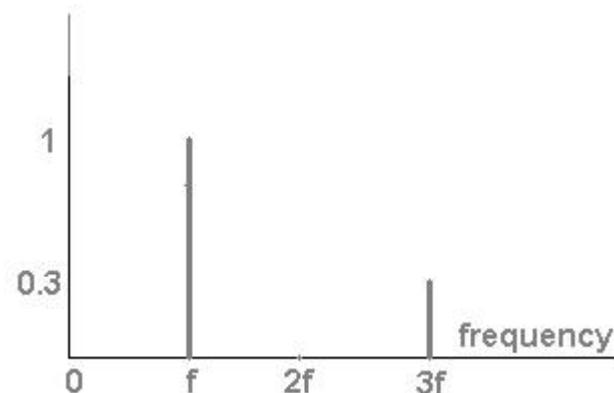
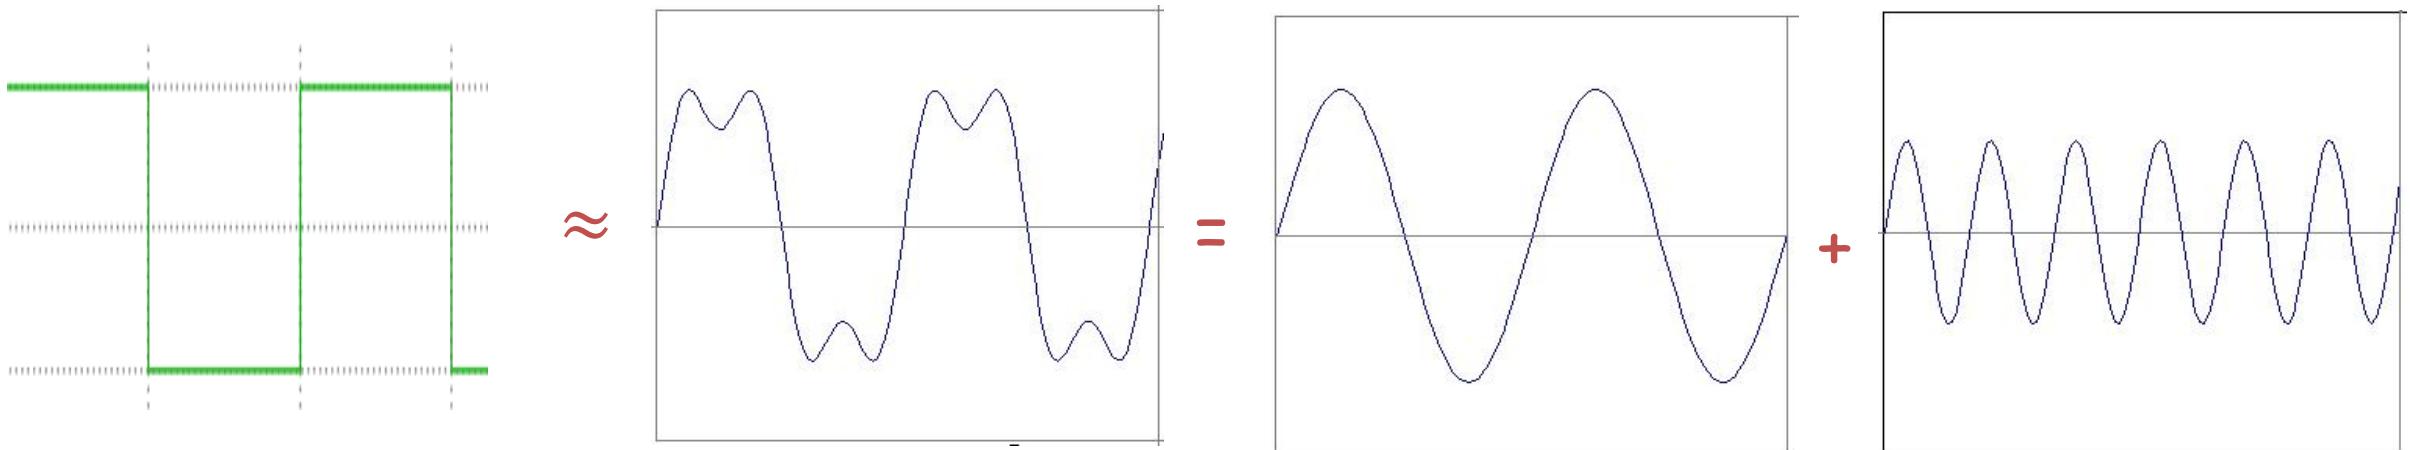
- Add enough of them to get **any** signal $g(x)$ you want!





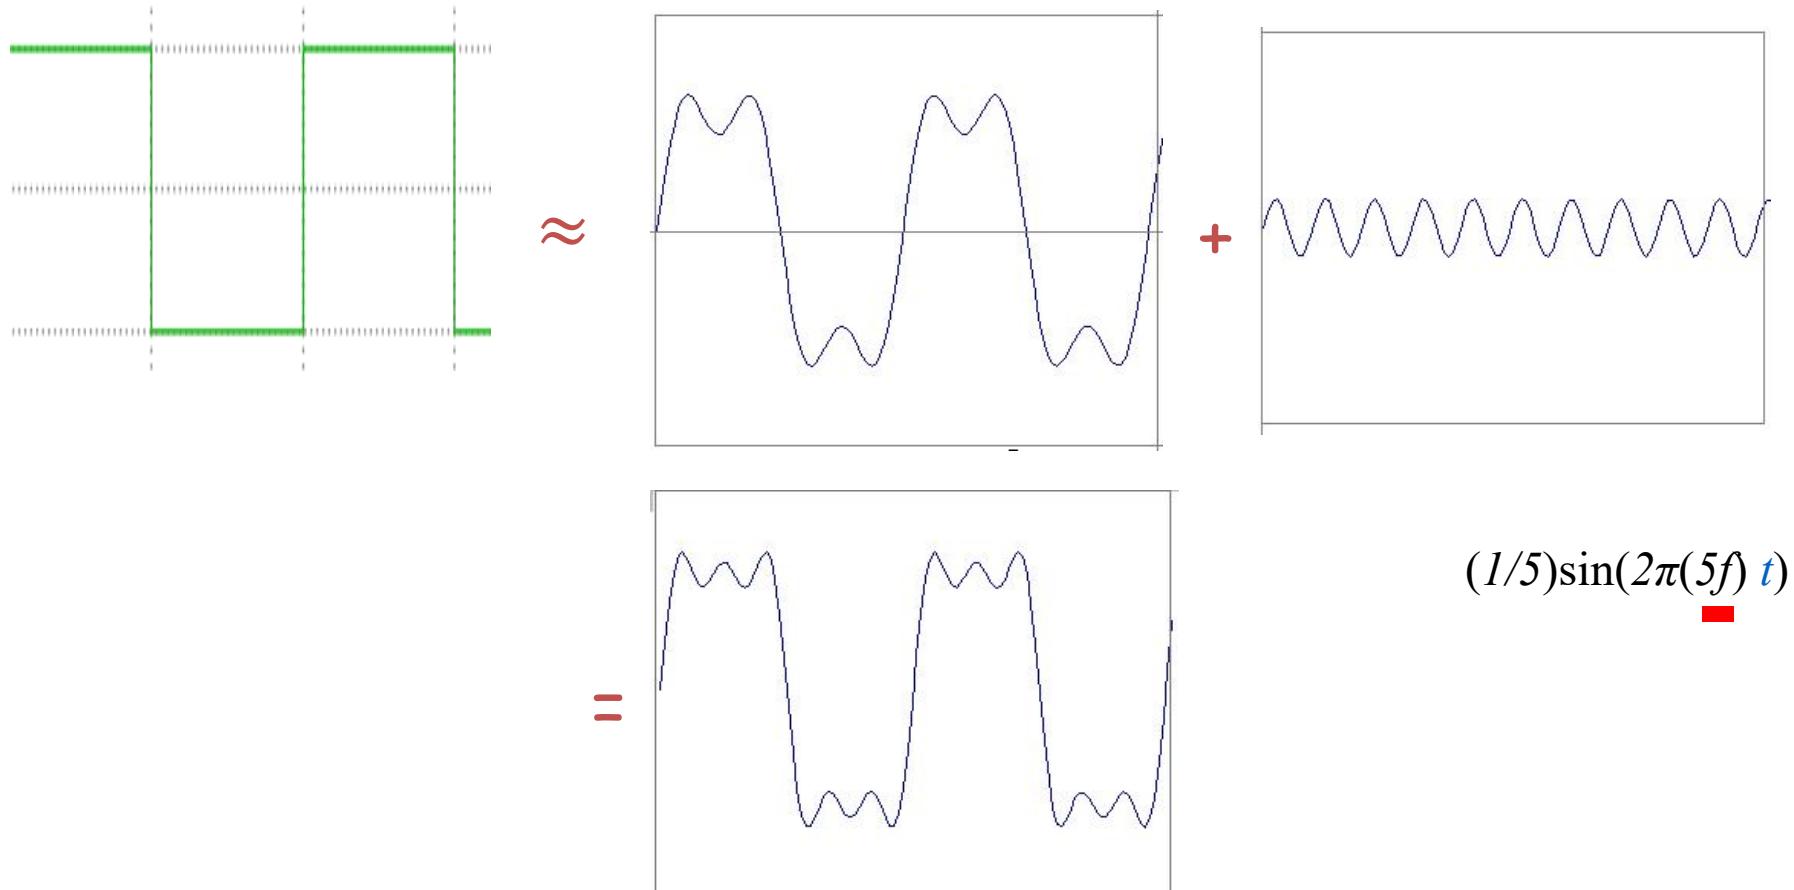
Frequency Spectra

- Example: $g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f)t)$



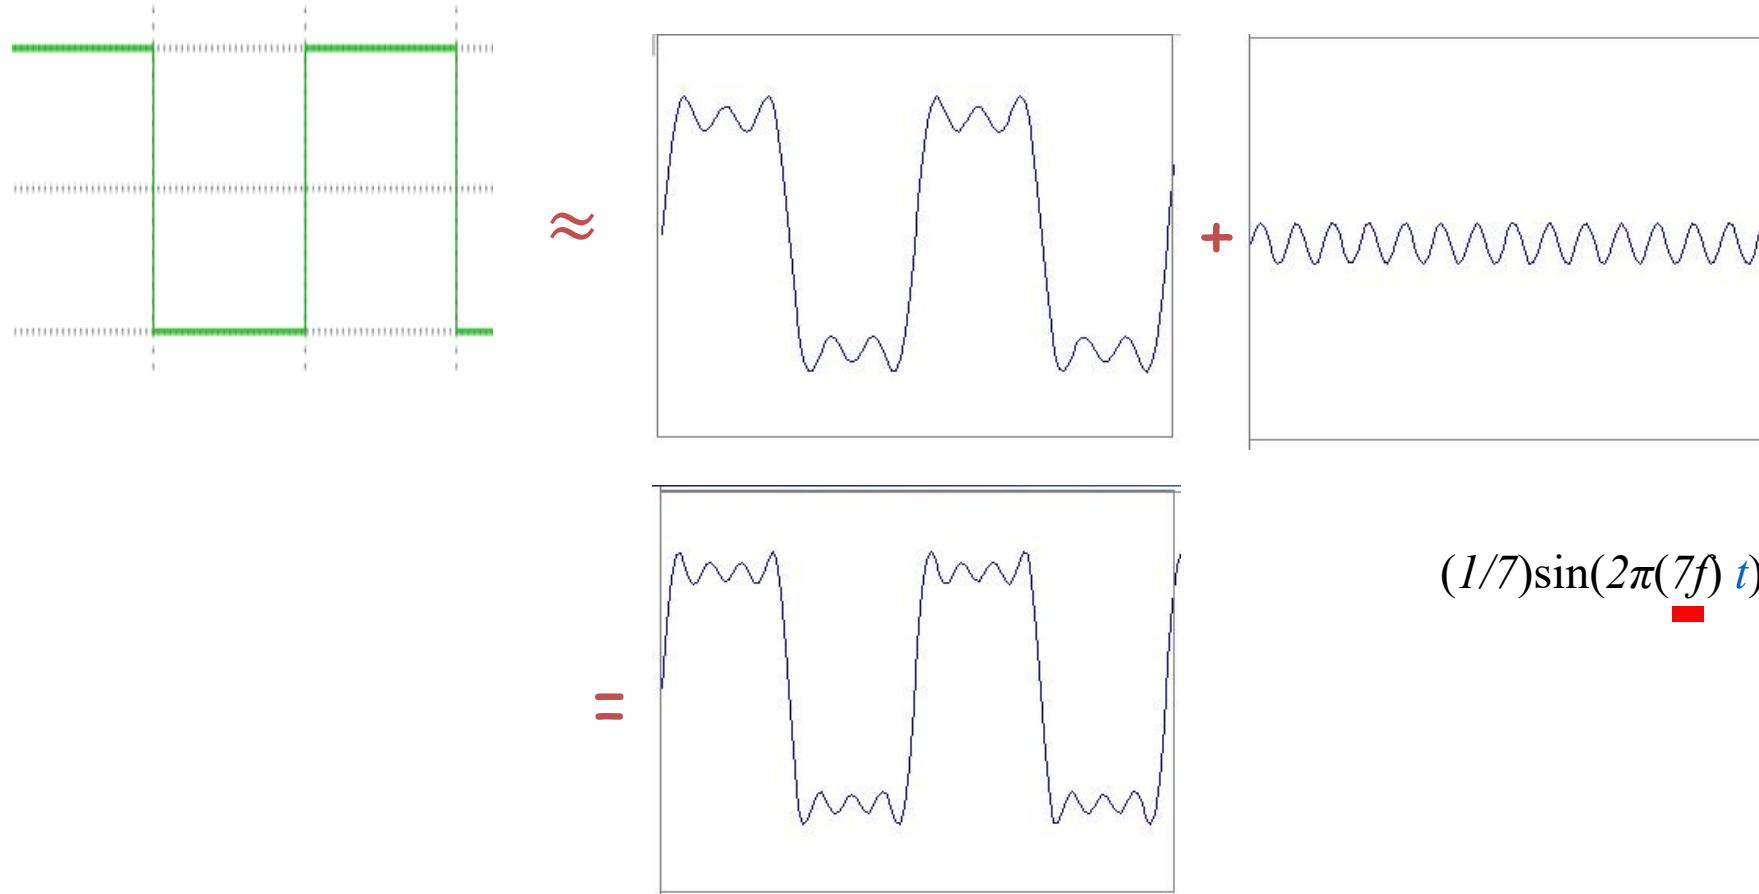


Frequency Spectra



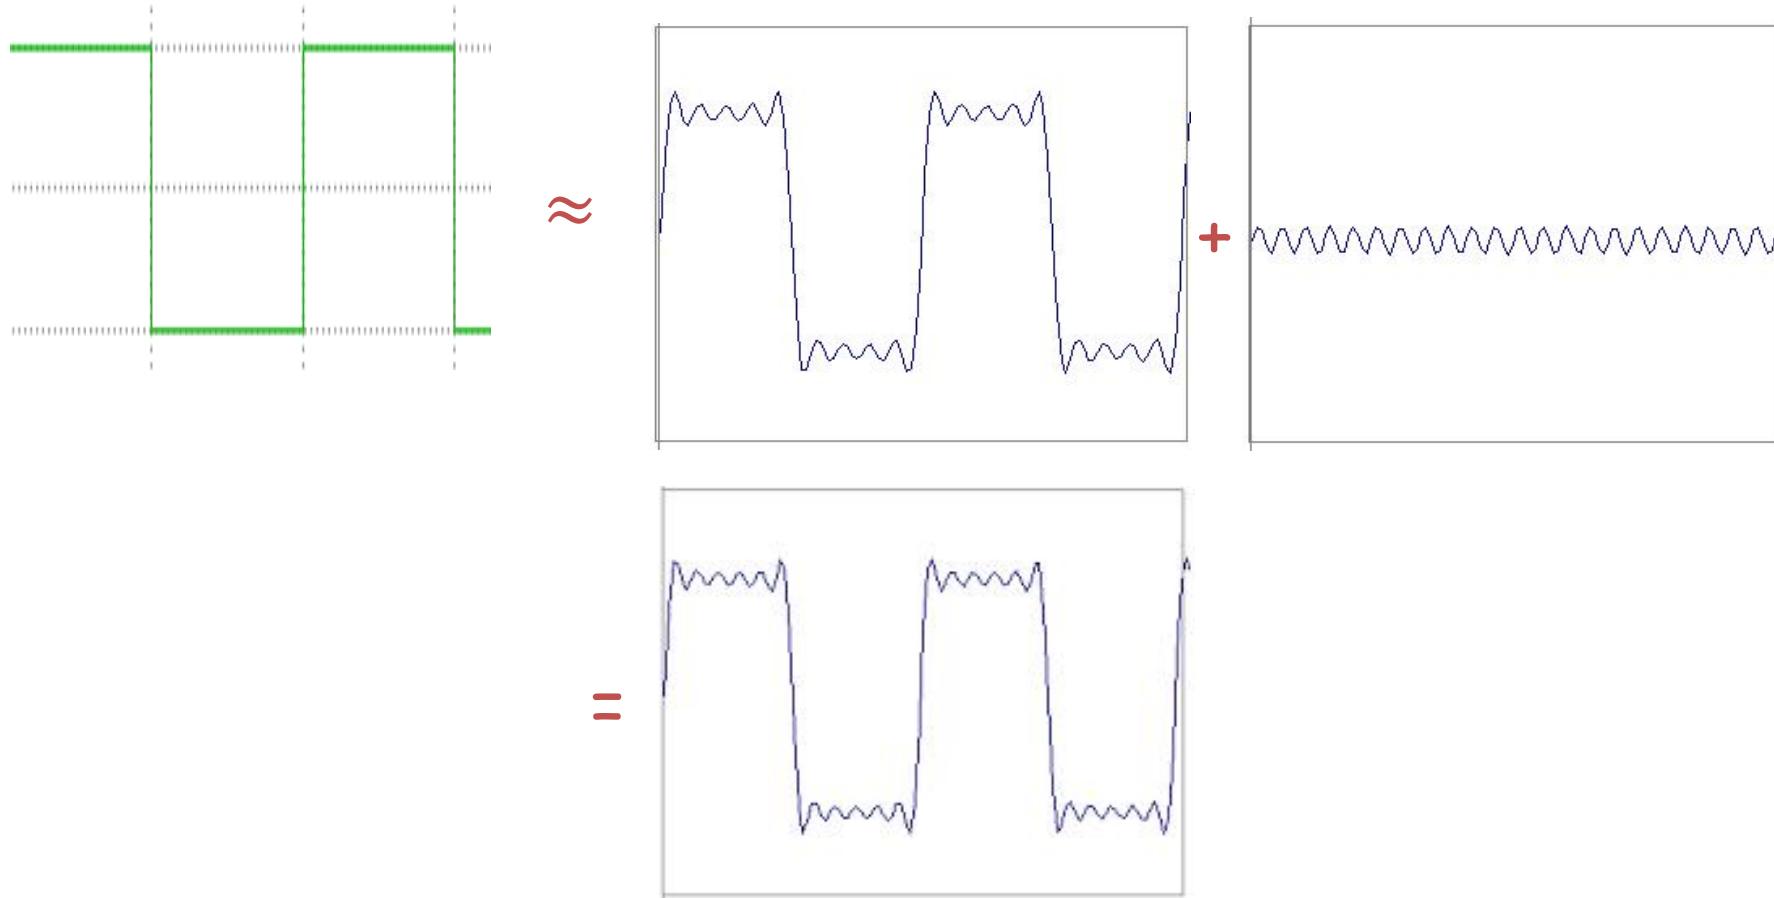


Frequency Spectra



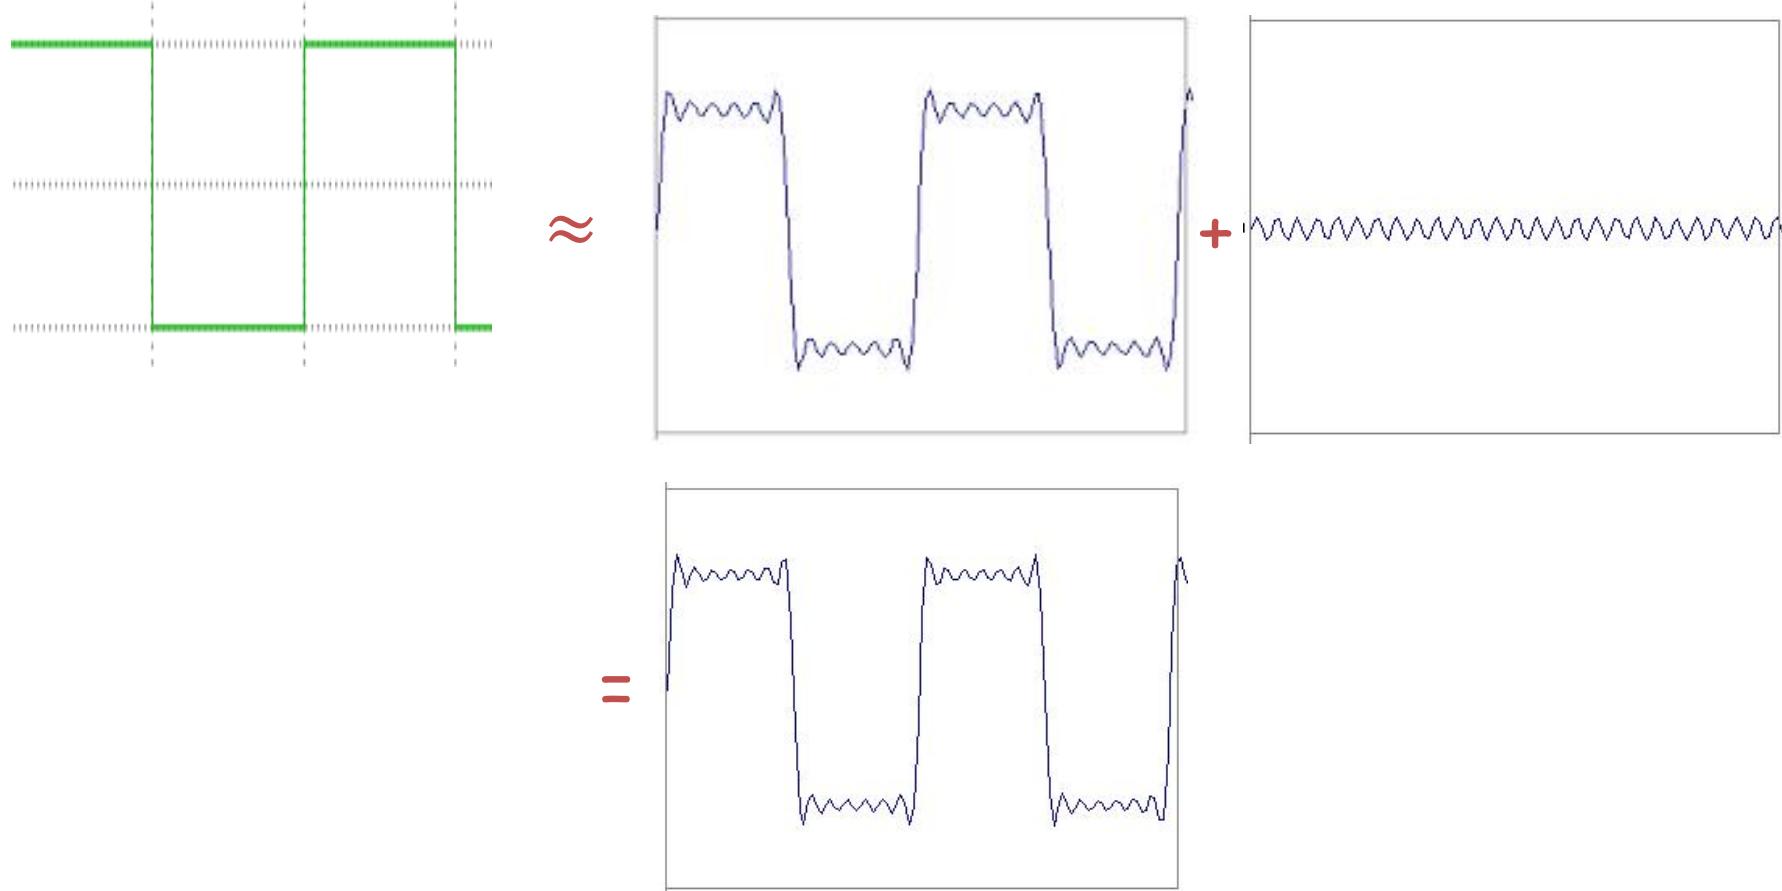


Frequency Spectra





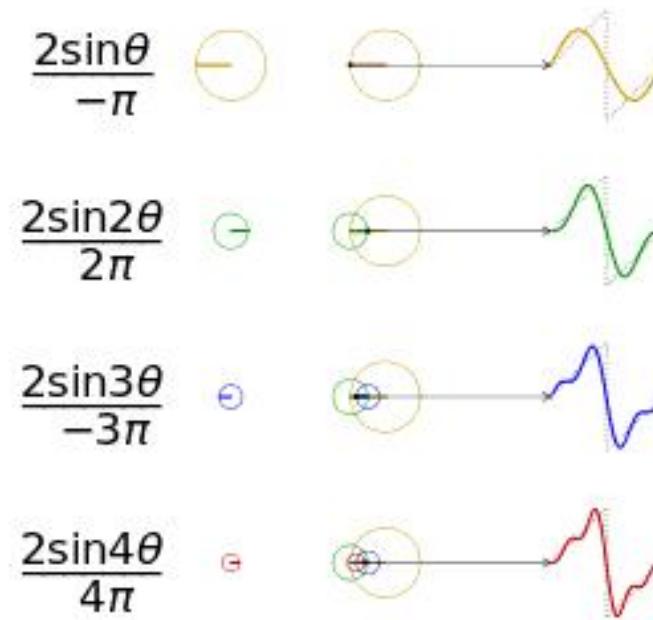
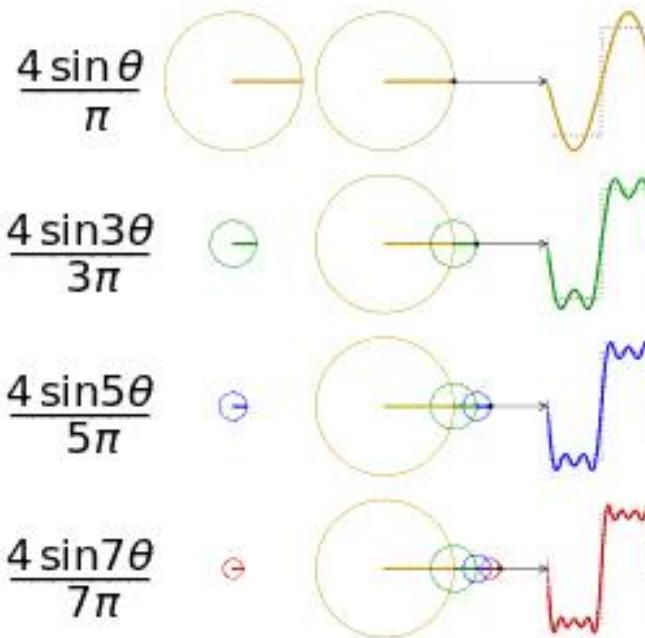
Frequency Spectra





Frequency Spectra

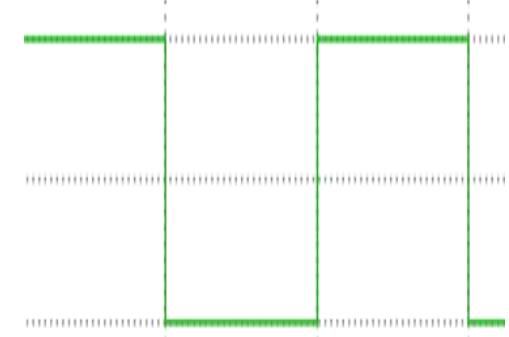
- Interesting explanations



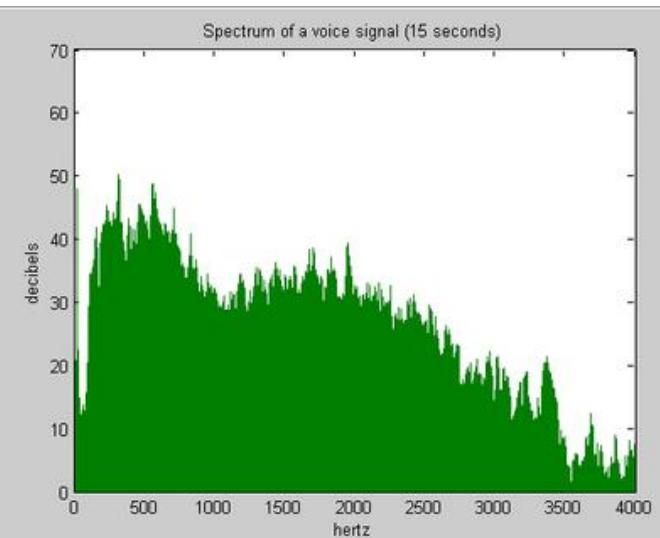
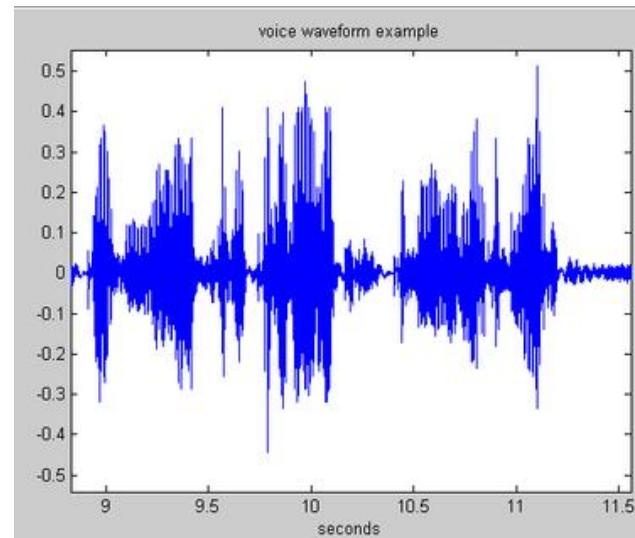
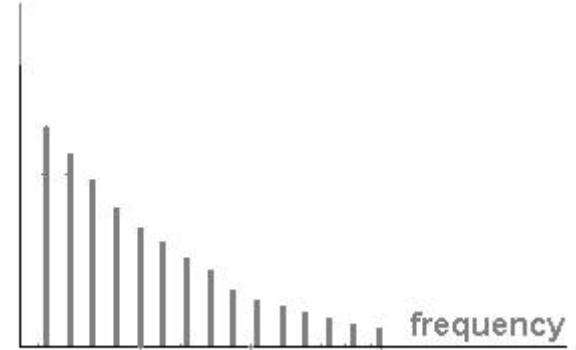


Frequency Spectra

- We think of music in terms of frequencies at different magnitudes



$$= A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi k t)$$





Two-dimensional Fourier Transforms

- An **oriented sinusoid** $s(x, y) = \sin(\omega_x x + \omega_y y)$
- The corresponding two-dimensional Fourier transforms

$$H(\omega_x, \omega_y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} h(x, y) e^{-j(\omega_x x + \omega_y y)} dx dy$$

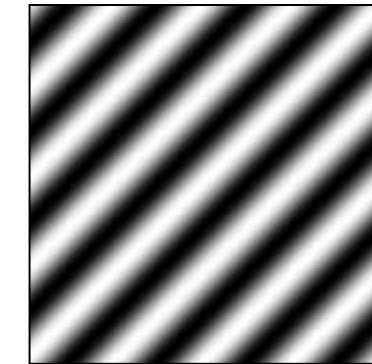
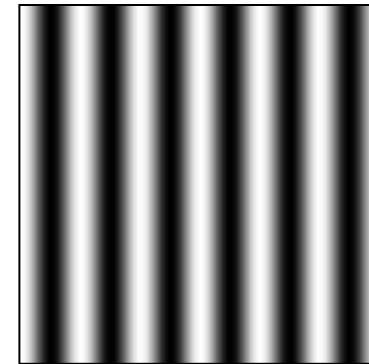
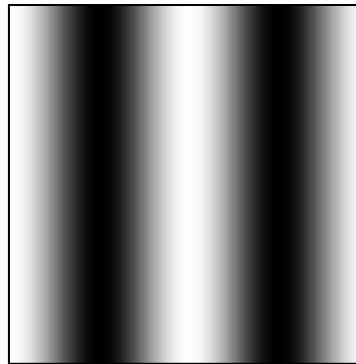
- In the discrete domain

$$H(k_x, k_y) = \frac{1}{MN} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} h(x, y) e^{-j2\pi \frac{k_x x + k_y y}{MN}}$$

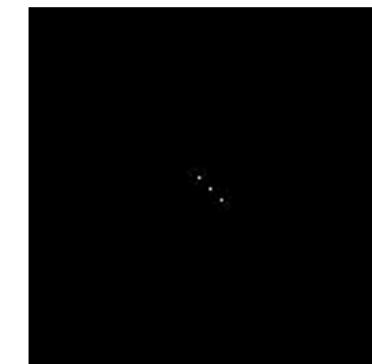
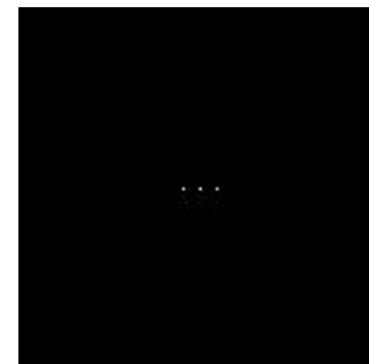
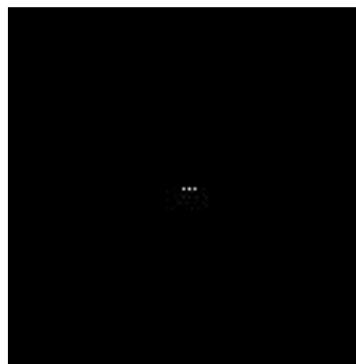


Fourier analysis in images

Intensity Image

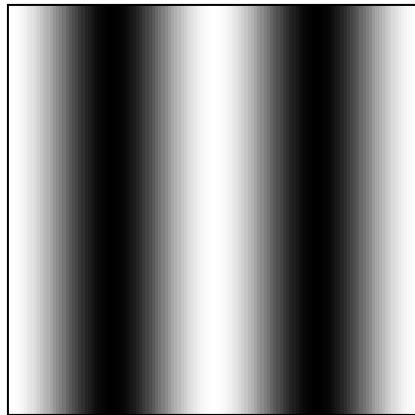


Fourier Image

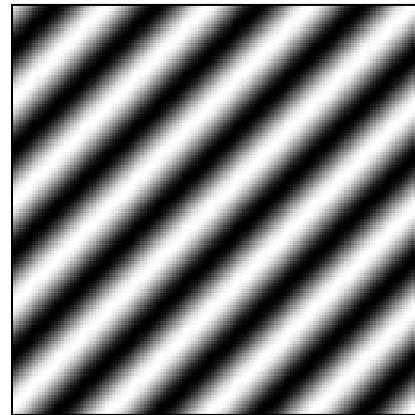




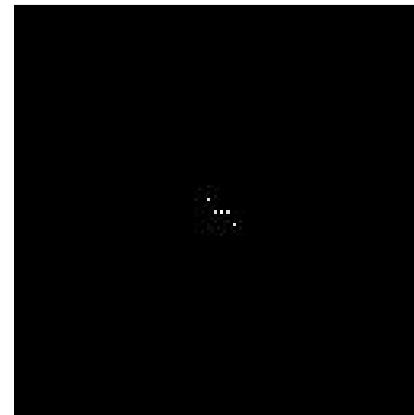
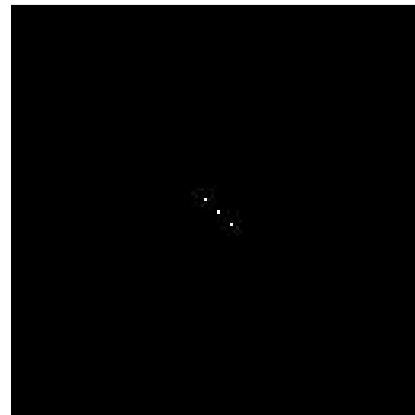
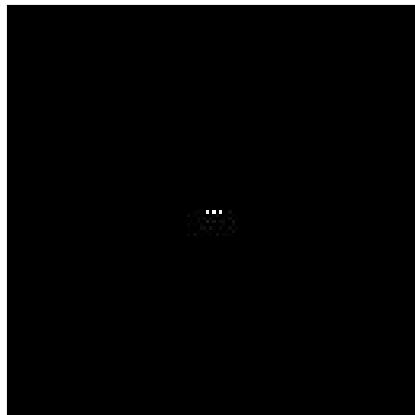
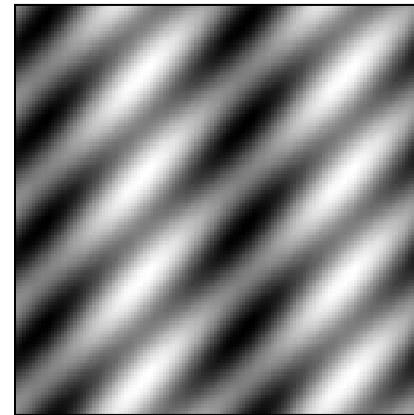
Signals can be composed: linear operation



+



=



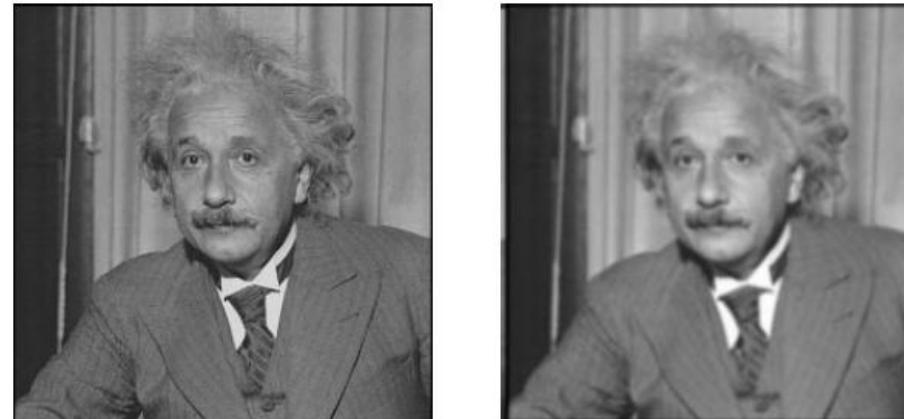


Example of 2D Fourier Transform

- **Smoothing**

二维傅里叶变换：

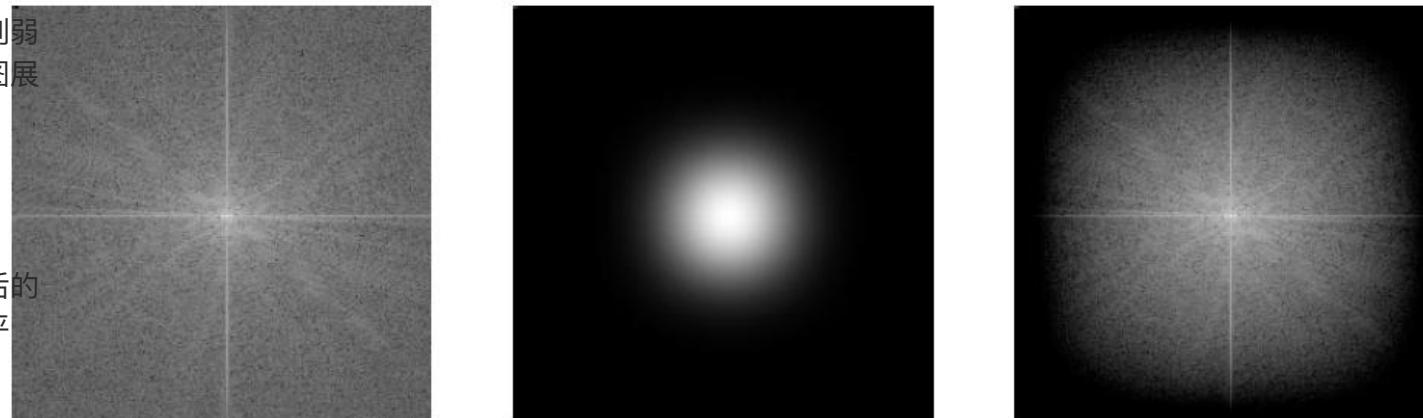
首先，图像被转化到频域。频域图像通常包含不同频率成分的信息。频率低的成分对应于图像的平滑区域，而高频成分则对应于细节和边缘。



在频域中，图像的高频成分（细节部分）被削弱或去除，从而实现了平滑处理。下方的频谱图展示了平滑处理后频率成分的变化。

结果：

原始图像在频域中有高频和低频成分，处理后的图像通过减少高频成分，使图像看起来更加平滑，细节减少。





在图像处理中，傅里叶变换常用于图像平滑（或去噪）

Fourier Transform

在图像去噪中，傅里叶变换可以将图像从时域转换到频域。在频域中，噪声通常表现为高频成分，而低频成分则代表图像的主要结构。

通过滤波（例如使用高通或低通滤波器）可以在频域中去除这些噪声成分。常见的做法是通过低通滤波器去掉高频部分，然后将其逆傅里叶变换回时域，从而实现去噪。

- Fourier transform stores the **magnitude** and **phase** at each frequency
 - Magnitude encodes **how much signal** there is at a particular frequency
 - Phase encodes **spatial information** (indirectly)
 - For mathematical convenience, this is often notated in terms of **real and complex numbers**
 - Fourier transform of a real signal is **symmetric about the origin**
- The Convolution Theorem

$$\mathcal{F}[g * h] = \mathcal{F}[g]\mathcal{F}[h]$$

$$g * h = \mathcal{F}^{-1}[\mathcal{F}[g]\mathcal{F}[h]]$$

在时域中对两个信号 g 和 h 进行卷积操作，等价于在频域中对这两个信号的傅里叶变换进行乘法操作

Pyramids



Pyramids

- What we have studied
 - All of the image transformations produce output images of **the same size** as the inputs
- **Change** the resolution of an image
 - Example: finding a face in an image
 - ✓ Accelerating the search for an object
 - ✓ Performing multi-scale editing operations
- Operations
 - Upsampling
 - Downsampling



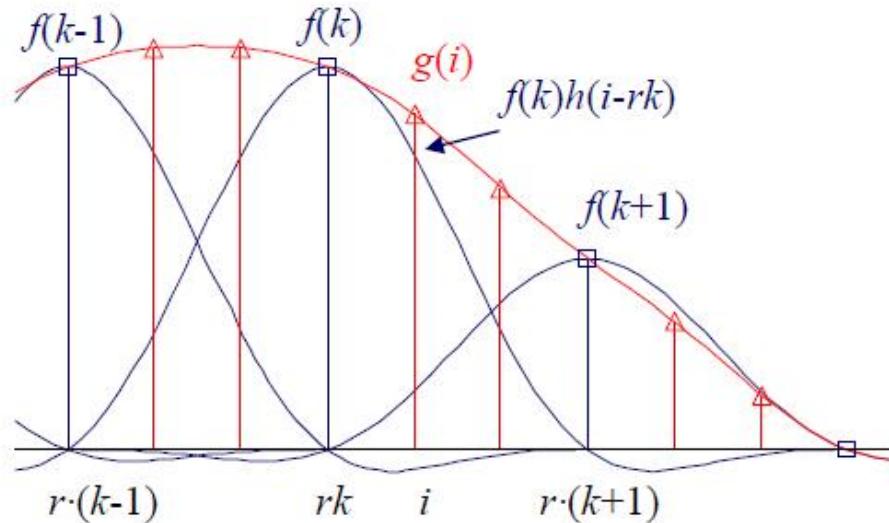
Interpolation

- Select some interpolation kernels to **convolve** the image

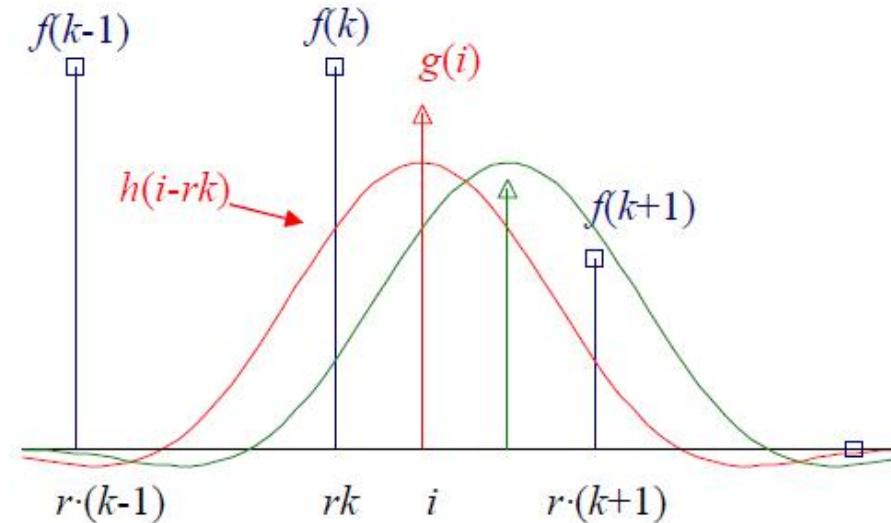
$$g(i, j) = \sum_{k,l} f(k, l) h(i - rk, j - rl)$$

$$g(i) = \sum_k f(k) h(i - rk)$$

the upsampling rate



weighted summation of input values



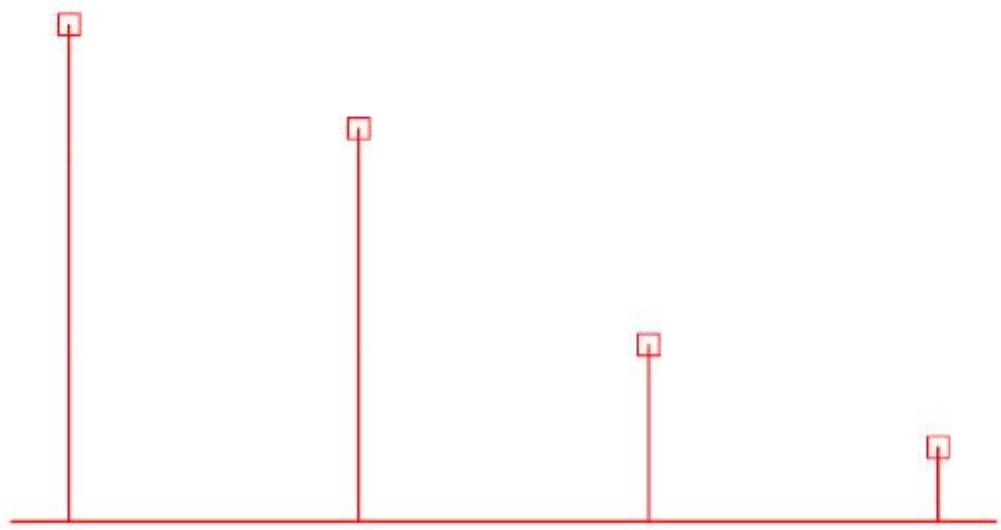
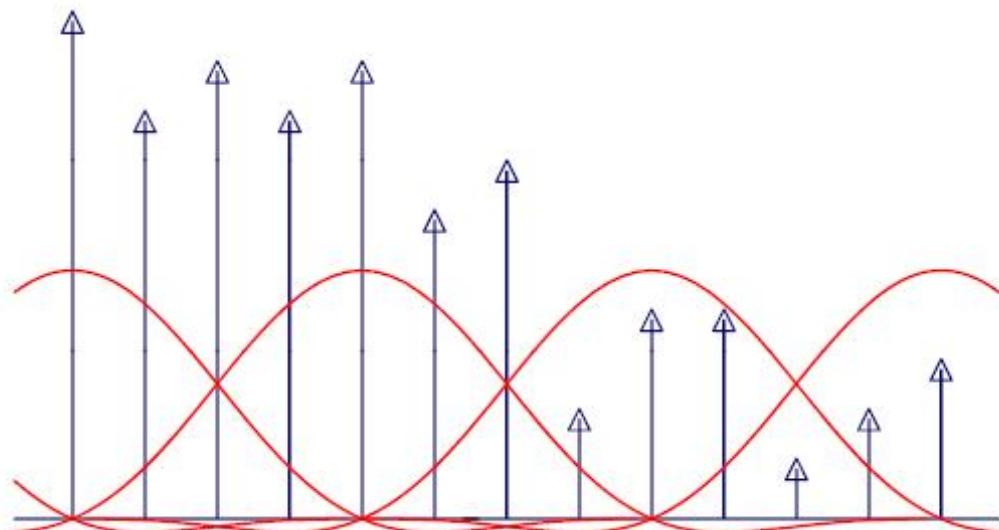
polyphase filter interpretation



Decimation

- Decimation (down-sampling) is required to reduce the resolution
 - Evaluate the **convolution** at **every r th sample (stride)**
 - ✓ First convolve the image with a low-pass filter
 - ✓ Then keep every r th sample

$$g(i, j) = \sum_{k, l} f(k, l)h(ri - k, rj - l)$$

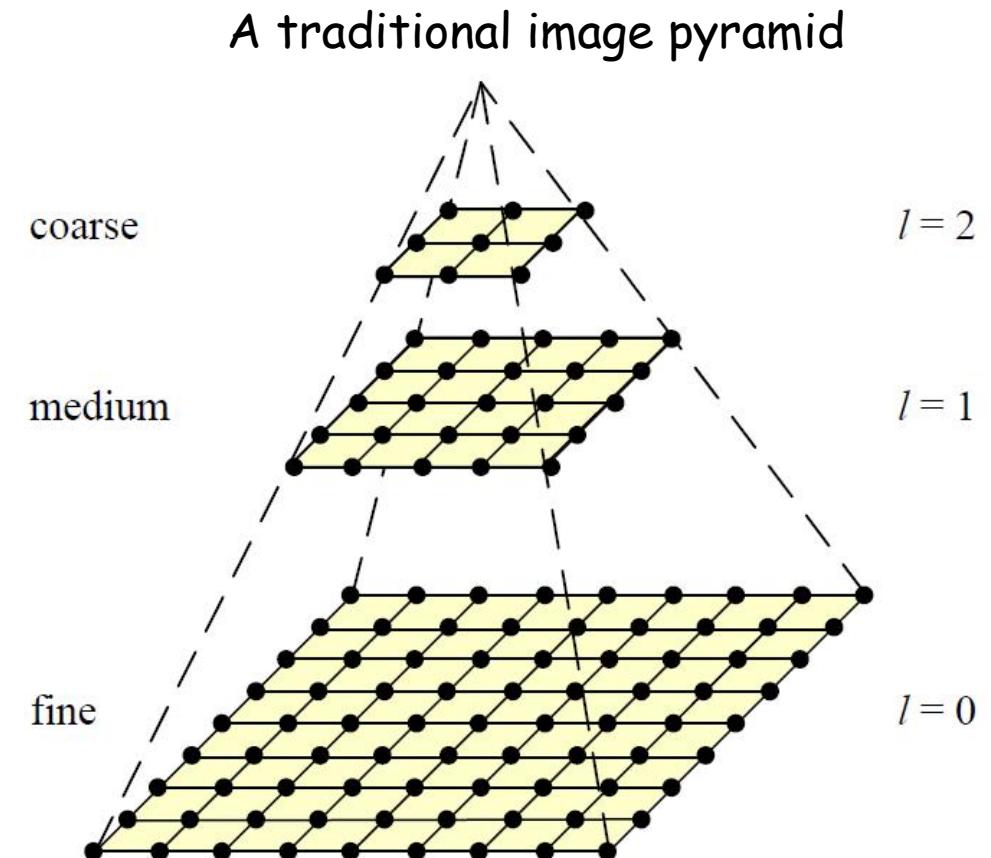




Multi-Resolution Representations

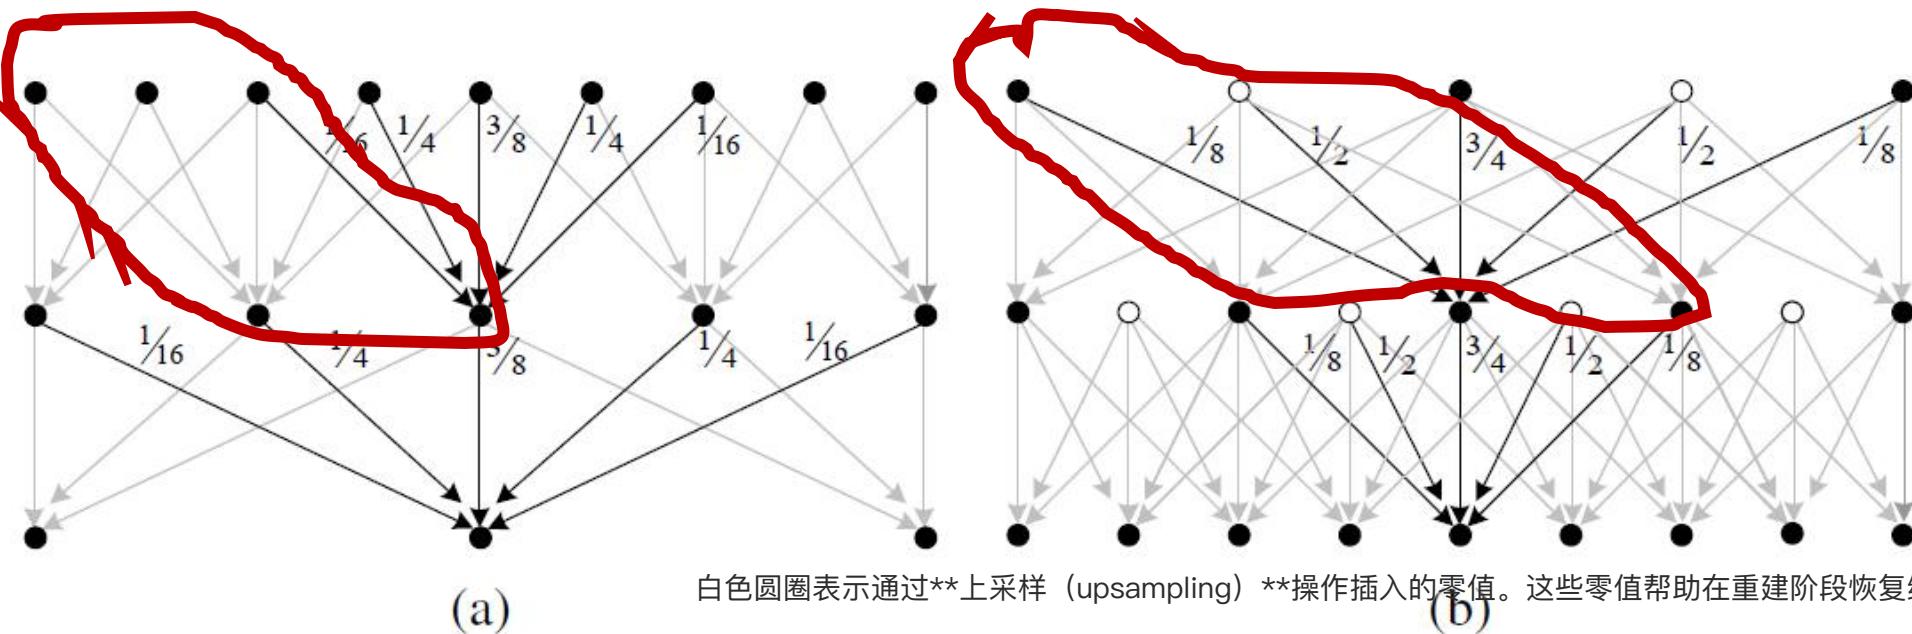
- Octave pyramids
 - Can be used to accelerate **coarse-to-fine** search algorithms

$$\frac{1}{16} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$





Multi-Resolution Representations



重建系数是分析系数的两倍，这意味着在图像重建时，计算需要使用更大的滤波器系数来确保图像细节的恢复。

Figure 3.33 The Gaussian pyramid shown as a signal processing diagram: The (a) analysis and (b) re-synthesis stages are shown as using similar computations. The white circles indicate zero values inserted by the $\uparrow 2$ upsampling operation. Notice how the reconstruction filter coefficients are twice the analysis coefficients. The computation is shown as flowing down the page, regardless of whether we are going from coarse to fine or *vice versa*.



Gaussian Pyramid

- **A Laplacian of Gaussian (LoG)** 描述了图像的变化率

$$\text{LoG}\{I; \sigma\} = \nabla^2(G_\sigma * I) = (\nabla^2 G_\sigma) * I,$$

$$\nabla^2 = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}$$

- **Difference of Gaussians (DoG)**
 - To obtain an approximation of the Laplacian of Gaussian

$$\text{DoG}\{I; \sigma_1, \sigma_2\} = G_{\sigma_1} * I - G_{\sigma_2} * I = (G_{\sigma_1} - G_{\sigma_2}) * I$$

它通过计算两个不同标准差的高斯滤波器的差异来近似Laplacian of Gaussian

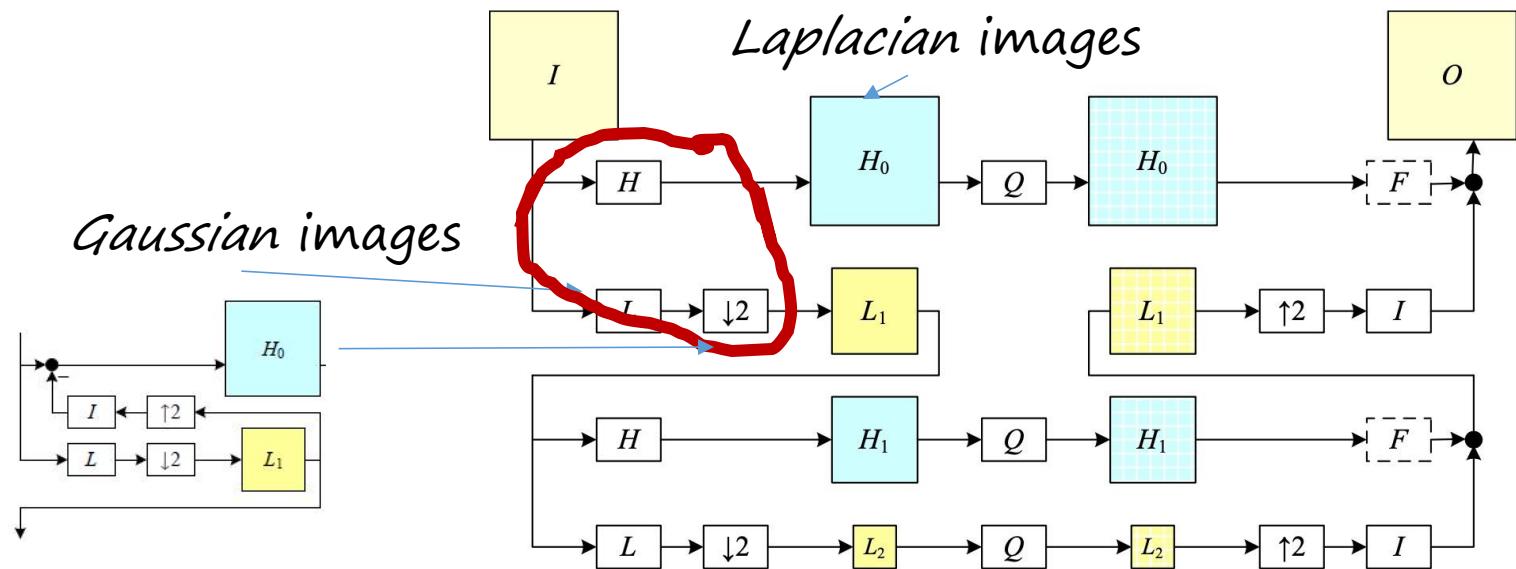


Multi-Resolution Representations

- To compute the Laplacian pyramid

➤ First interpolate a lower resolution image to obtain a reconstructed **low-pass (short-pass) version (high frequency)**

➤ Then **subtract this low-pass version from the original to yield the band-pass "Laplacian" image**



1. 拉普拉斯金字塔 (Laplacian Pyramid)
的计算:
计算方法:

插值低分辨率图像: 首先, 我们从一个低分辨率图像开始, 使用插值技术生成它的高频部分。高频部分通常是图像的细节部分。

减去低通版本: 然后, 通过从原始图像中减去该低分辨率的“低通”版本 (即平滑版本), 我们可以得到**带通 (band-pass)**的拉普拉斯图像, 它包含了图像的细节信息。

原理:

高频信息是图像中最细致的部分, 比如边缘和小细节。

低频信息是图像的整体结构和模糊部分, 描述了图像的较大区域的变化。

Why? THINKING LOG and DOG
拉普拉斯金字塔通过从原始图像中提取出这些高频频分, 可以帮助我们在多分辨率表示下更清楚地分析图像细节。

图示:

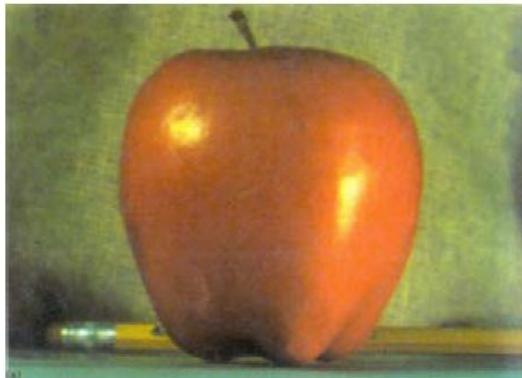
图中的流程展示了如何从高斯金字塔 (Gaussian pyramid) 中提取低频信息, 并通过差值操作得到高频的拉普拉斯图像。此操作可以帮助我们更好地分析和重建图像的细节部分。
What looks like? DL



Application: Image Blending

- Keys:

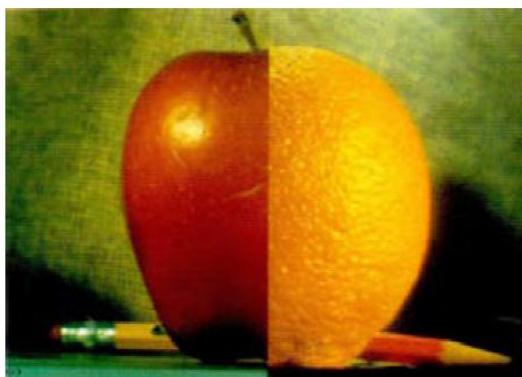
- The low-frequency color variations between the red apple and the orange are **smoothly blended**
- The higher-frequency textures on each fruit are **blended more quickly** to avoid “ghosting” effects when two textures are overlaid



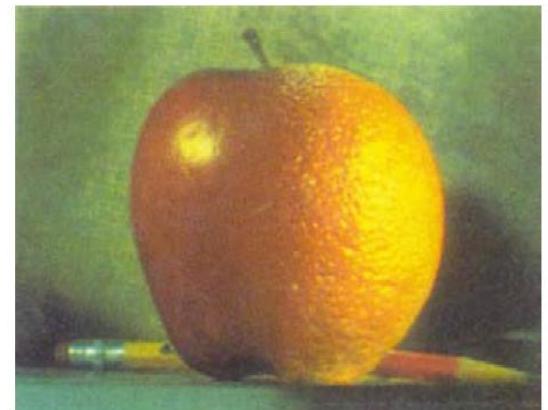
(a)



(b)



(c)



(d)

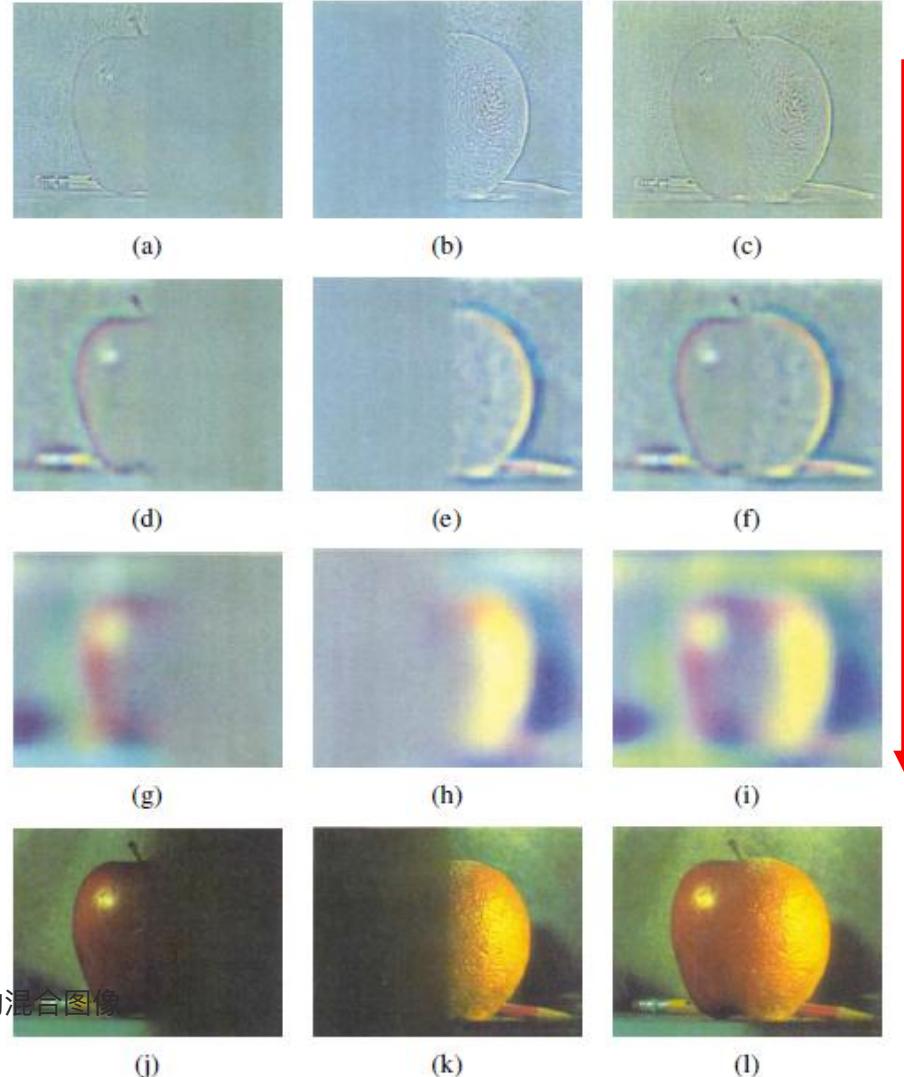
(c) regular splice, (d) pyramid blend.



Application: Image Blending

- Steps:

- Each source image is first decomposed into its own **Laplacian pyramid**
- Each band is then **multiplied by a smooth weighting function**
 - ✓ Create these weights is to take a binary mask image
- Each Laplacian pyramid image is then multiplied by its corresponding **Gaussian mask**
- The sum of these two weighted pyramids is then used to construct the final image



将两张图像在不同的频率层上分别加权融合（使用平滑权重掩膜），再重构出自然、平滑过渡的混合图像

Geometric Transformations

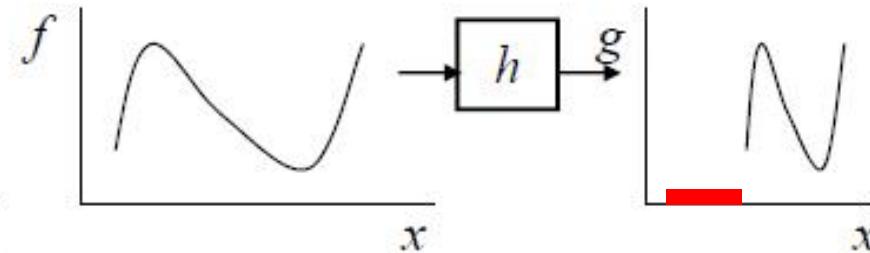
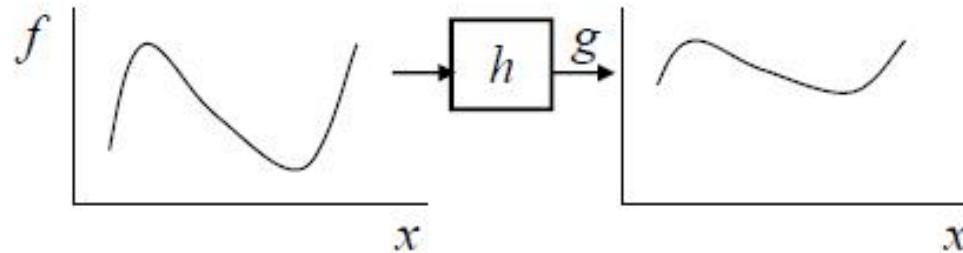
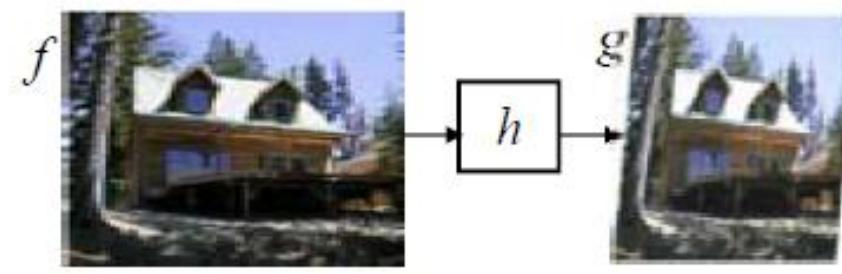
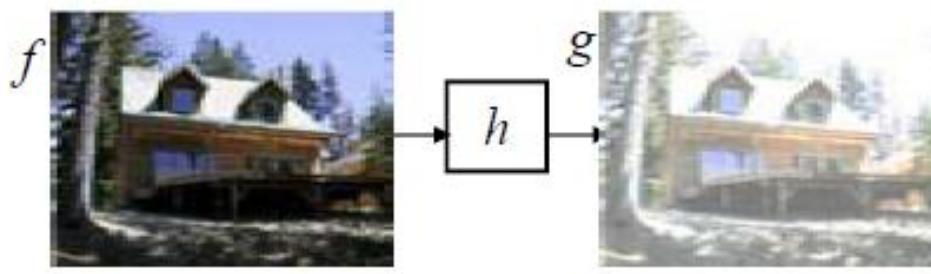


Geometric Transformations

- General transformations

- Image rotations
- General warps

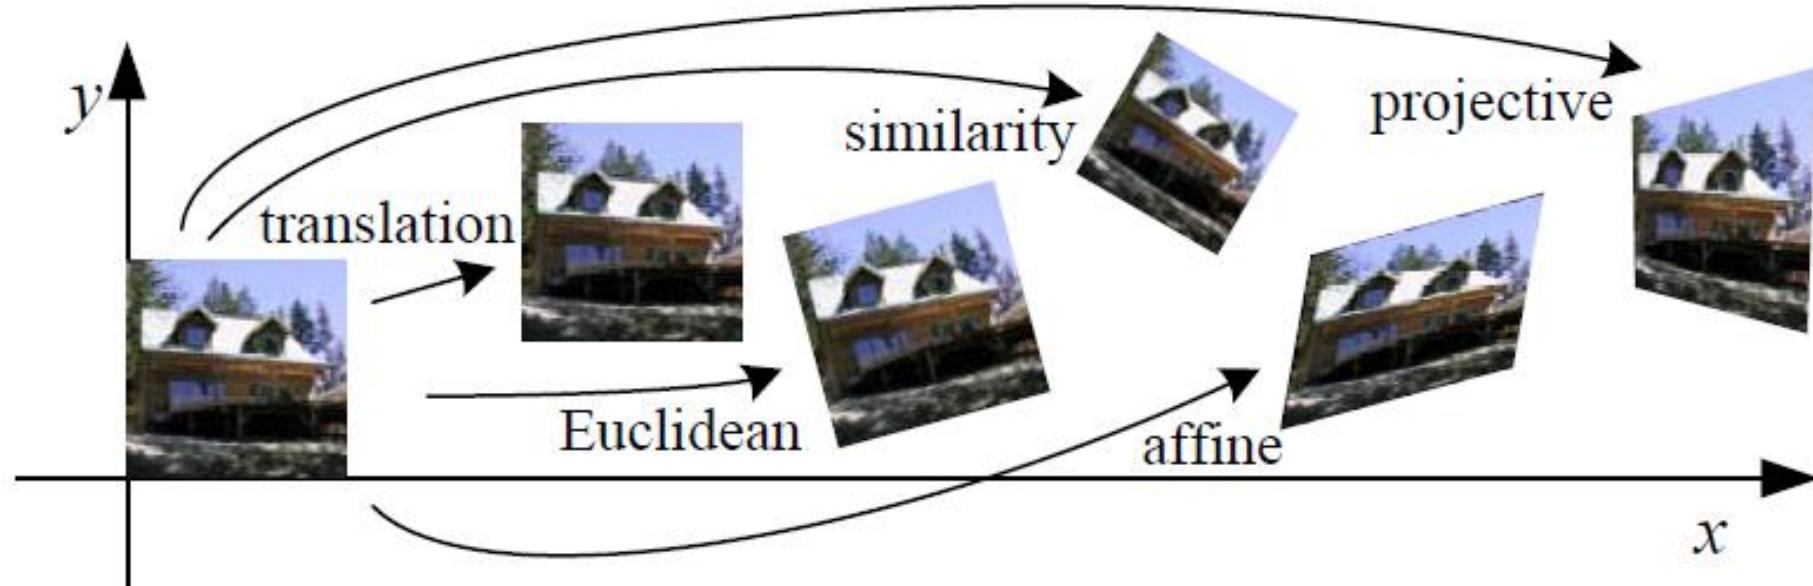
$$g(\mathbf{x}) = h(f(\mathbf{x}))$$





Geometric Transformations

- Basic set of 2D geometric image transformations





Parametric Transformations

- Apply a **global** deformation to an image
- Controlled by a **small number** of parameters

Transformation	Matrix	# DoF	Preserves	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	



Forward Warping Algorithm

- Limitations:
 - Target location has a **non-integer value**
 - ✓ round the value
 - ✓ "distribute" the value
 - The second major problem with forward warping is the appearance of **cracks and holes**

procedure *forwardWarp*(*f*, *h*, **out** *g*):

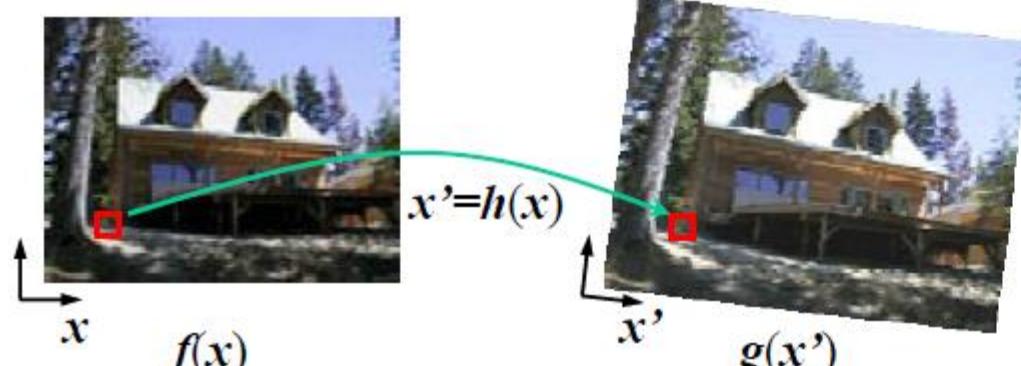
For every pixel x in $f(x)$

1. Compute the destination location $x' = h(x)$.
2. Copy the pixel $f(x)$ to $g(x')$.

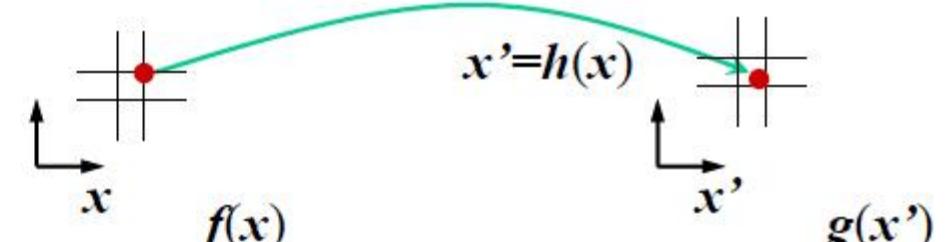


Forward Warping Algorithm

- An example



(a)



(b)



Inverse Warping

- Resampling an image at **non-integer locations** is a well-studied problem
- **High-quality filters** that control aliasing can be used

procedure *inverseWarp*(f , h , **out** g):

For every pixel x' in $g(x')$

1. Compute the source location $x = \hat{h}(x')$
2. Resample $f(x)$ at location x and copy to $g(x')$



Mesh-based Warping

- Local deformations with more degrees of freedom are often required
 - Changing the appearance of a face from a frown to a smile
 - Different amounts of motion are required in different parts of the image
- First approach
 - Specify a sparse set of corresponding points
 - Displacement of these points can then be interpolated to a dense displacement field
 - Alternative methods for interpolating a sparse set of displacements include moving nearby quadrilateral mesh vertices



Mesh-based Warping

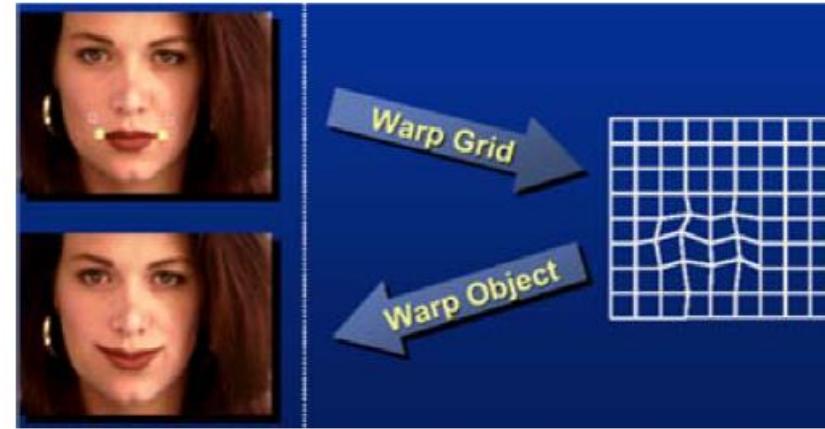
- A second approach
 - Specifying displacements for local deformations is to use **corresponding oriented line**
 - ✓ Pixels along each line segment are transferred from source to destination **exactly** as specified
 - ✓ Other pixels are warped using a **smooth interpolation** of these displacements
 - ✓ Each line segment correspondence **specifies** a translation, rotation, and scaling
- One possibility for specifying displacement fields is to **use a mesh specifically** adapted to the underlying image content



Mesh-based Warping

- Examples

- (a) sparse control **points**->deformation grid;
- (b) denser set of control point correspondences;
- (c) oriented **line** correspondences;
- (d) uniform quadrilateral grid.



(a)



(b)



(c)



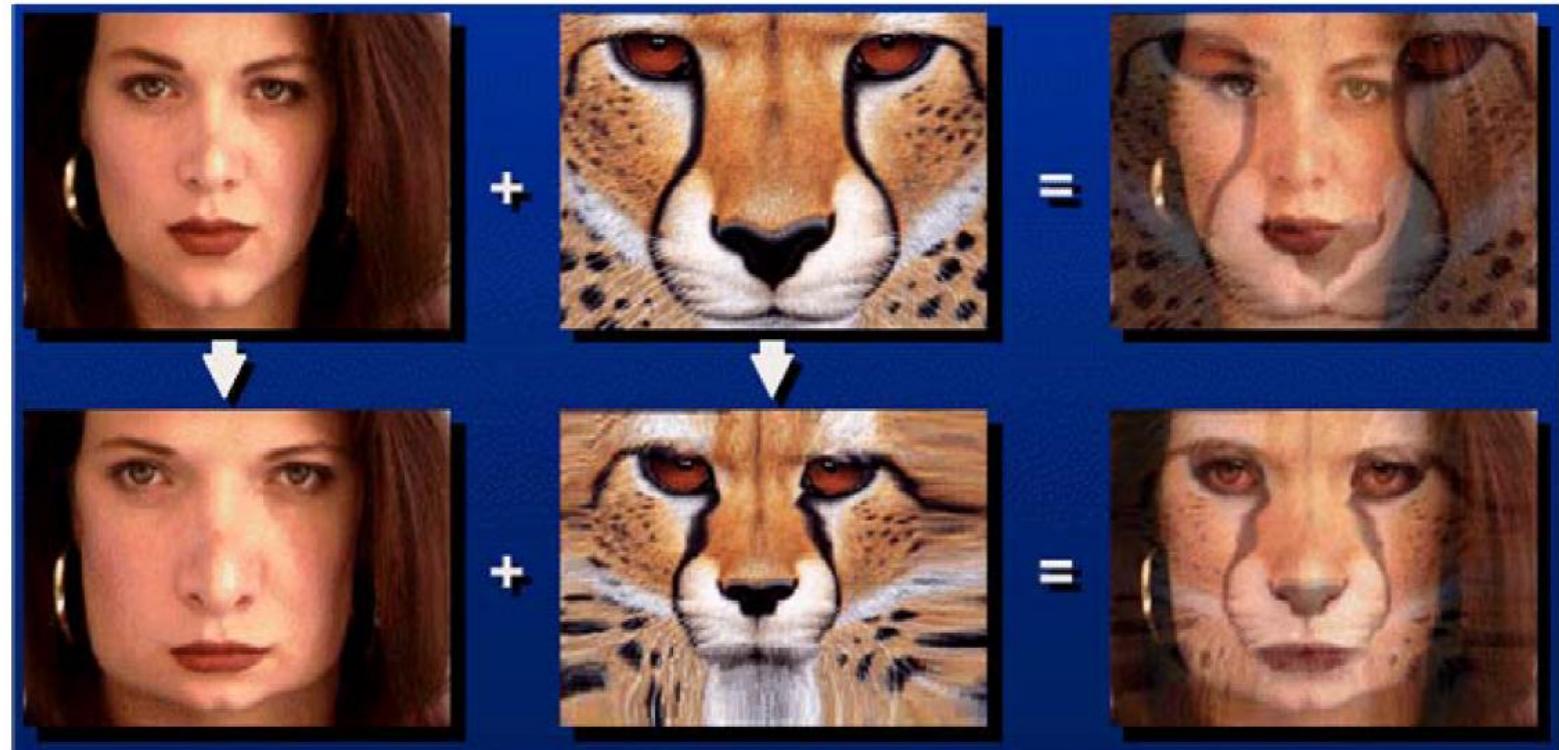
(d)



Image Morphing

- Top row: if the two images are **just blended**, visible ghosting results
- Both images are first warped to the same **intermediate location**

变形到同一中间位置

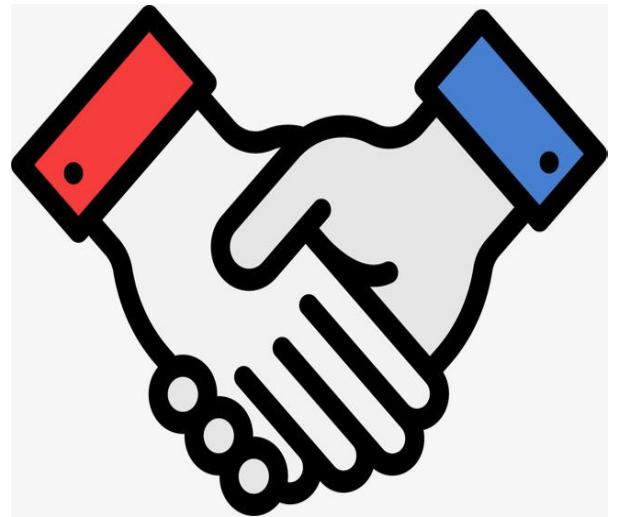


Conclusions



Conclusions

- Thinking in Frequency
 - Remove noise (high frequency)
 - Theory
- Pyramids
 - Upsampling
 - Downsampling
- Geometric Transformations
 - Inverse warping (augmentation)
 - Mesh-based warping



Thanks



zhengf@sustc.edu.cn